


Solution 1.

Figure 1: Counterexample for MODIFIEDPRIM

A counter example is provided as above.

Figure 1.1 is the graph to be traversed.

By the provided algorithm, all the vertices in the tree T will be full in the order of edge weight: edge(A,B), edge(A,C) (as indicated in line 4 for the set S). We then end up with Figure 1.2. The vertex D is thus not considered and the algorithm returns **infeasible**.

However, we can see clearly that a MST is indeed feasible, as in Figure 1.3. Thus the algorithm is not correct in that it returns **infeasible** when in fact an MST is feasible.

Solution 2.**(a)****High-level understanding**

We do three iterations. **The first iteration** takes $O(ml)$ which traverses through every single bus stop to create a graph G . In this graph G , every bus line B_i is a vertex and edges represent the connection between two bus lines. **The second iteration** takes $O(m^2)$ time. It produces every possible option represented by an array $[B_1, B_2, \dots, B_n]$ meaning boarding B_1 then changing into B_2 and so on. The travel option array is added as a vertex attribute to the corresponding vertex. It also calculates the bus fare incurred by it successively and add it as a vertex attribute. **The final iteration**, which takes $O(m^3)$ time, considers all the travel options produced by the second iteration which ends at the destination stop b and sort them by the bus fare. Then it iterates from the cheapest travelling option to the more expensive ones. For each option, it checks if it satisfies the time constraints. If it does, the algorithm stops immediately and export that option in the correct format.

Detailed description

get-graph() To convert it into a graph, we traverse from a to find a set S_0 of all the bus lines a is on. Then for each bus line B_i that we just found than go through a , we traverse through all the possible bus stops in that bus line. If that bus stop is also on some other bus lines other than the ones in S_0 , then we now have an edge between the current bus line that the new bus line. We then go on and traverse through all the bus stops in the new bus line and repeat the process.

get-all-options() We first add all the bus lines B_i that go through a to a set A , and the ones go through b to a set B . The set of vertices A serves as a set containing all vertices that have been visited. We have an array stored as an attribute of the vertex initialized with all vertices in A , e.g., $[B_1], [B_2], \dots, [B_n]$. Then for each vertex (e.g. B_k) (bus line) in G but not in A that is connected to any vertex in A , we attach that to the corresponding array $[B_2, B_k]$ where we can transfer from B_2 to B_k . We also add them to the set A as visited vertices.

Eventually, we do this for every vertex that is yet to be visited $v \in G, v \notin A$. Suppose v is adjacent to $B_i \in \{B_1, B_2, \dots, B_n\}$, we first add it to the visited vertices set A . Then, we look at all the arrays stored in B_i , e.g., $[B_1, B_4, B_2, B_i]$. All these arrays represent a path from a to the bus line represented by the last vertex. If it is not the same as any of the elements in the corresponding array (i.e. adding this bus line causes a loop), we add it to the array and store it in the new vertex v . We also calculate the bus fare for this renewed array option by adding the bus fare of v to the corresponding attribute of the previous vertex B_i . After adding, we add the bus fare for B_k on top of the price attribute of the previous vertex, e.g., B_2 , and add the updated price to B_k .

get-cheapest-option-in-time() We first iterate through all the bus options array and cheerypick the one ending with b destination stop. We first sort all the options by the bus fare. Then we iterate through all the options from the "cheapest to the most expensive ones. For each option, we check if it satisfies the time constraints. If it does, the algorithm stops immediately and export that option in the correct format.

How to check for time constraints? For each bus line, we have another time-pair array aside from the travel options array. We always add the the stop-time

pair of the bus stop to the time-pair array where we change bus line into the *travel plan array*. We go from the first bus line B_1 . We check if there is any buses with a departure time t later than t . Then we follow along to the intersecting bus stop to the second bus line B_2 and changes buses. We check if the time of getting off is earlier than the earliest available departure time of B_2 from that stop. We repeat the same process for every bus line B_i afterwards and always get on the next available bus to change into. Ultimately, we reach b and check if the arrival time is still earlier than t' . If so, then this travelling option is still valid and we stop and output the time-pair array.

(b) First, we prove the correctness of *get-graph()*.

Claim 2.1. *All instances where we can change from bus B_1 to B_2 is matched to exactly one edge between the corresponding vertices (bus lines).*

Proof. The *get-graph()* function builds a graph where each vertex represents a bus line, and an edge exists between two vertices if there is a bus stop that is common between the corresponding bus lines.

We consider two bus lines, B_1 and B_2 . We need to show that an edge exists between B_1 and B_2 in the graph if and only if there is a bus stop common to both B_1 and B_2 .

Let's consider a bus stop s which is common to both B_1 and B_2 . When the *get-graph()* function traverses through all the bus stops in B_1 , it will reach s and find that s is also on B_2 . Thus, it will add an edge between B_1 and B_2 in the graph. This proves the "if" part of our claim.

For the "only if" part, let's assume that an edge exists between B_1 and B_2 in the graph. This edge would have been added by the *get-graph()* function only when it found a bus stop that is common to both B_1 and B_2 . Thus, there must exist such a bus stop.

Therefore, we can conclude that all instances where we can change from bus B_1 to B_2 are matched to exactly one edge between the corresponding vertices (bus lines) in the graph. This validates the correctness of the *get-graph()* function. \square

Claim 2.2. *The *get-all-options()* function generates all possible sequences (without a loop) of bus lines from a to b and correctly calculates the total fare for each sequence.*

Proof. We first prove that it generates all possible sequences (without a loop) of bus lines from a to b . We prove by induction.

Base case We have only one bus line B_1 in the set A . Then we have B_1 in the array $[B_1]$.

Inductive hypothesis For all the visited vertices in the set A , the algorithm has generated all possible sequences (without a loop) of bus lines from a up to that visited vertex.

Inductive step For the current vertex A_i in A that have been visited, we look at all the adjacent vertices v of A_i . For each adjacent v that is not in the travel plan array which can incur a loop, we add it to that array. That means we have **exhausted** all our options for the next step for that A_i . Thus, we have generated all possible sequences (without a loop) of bus lines from a up to that visited vertex.

We then also prove by induction that the algorithm correctly calculates the total fare for each sequence.

Base case For the case where we only have one bus line B_1 in the set A , the total fare is the fare of B_1 . Thus the correctness is proven for this case.

Inductive hypothesis For each travel plan sequence, our algorithm correctly calculates the total fare.

Inductive step We calculate the bus fare for this renewed array option by adding the bus fare of v to the corresponding attribute of the previous vertex B_i . Since that is the only difference between the previous bus fare and the current bus fare, our algorithm is correct. \square

Claim 2.3. *The get-cheapest-option-in-time() function returns the cheapest travel plan that satisfies the time constraints.*

Proof. The correctness lies in whether or not we can correctly decide if it satisfy the time constraints.

We prove by induction.

Base case For a single bus line B_1 in the travel option array, we because our algorithm make sure the the departure time of B_1 from a is later than t , it is correct.

Inductive hypothesis Suppose that for i -th element in the array, the time constraints are satisfied.

Inductive step We only need to prove that for the new element, it departs later than the existing bus time and ends before the t' closing time. Because our algorithm checks that, the overall correctness is proven. \square

(c) *get-graph()* takes $O(ml)$ time. Since it goes through all possible bus stops and check if the intersecting bus lines have been added previously.

get-all-options() takes $O(m^2)$ time. Let us suppose that we have a_0 elements in the original set A and a_n elements in the original set B . Then we have a_1 vertices adjacent to any of the vertices in A and we add them to A . And a_2 vertices for the updated A and so on. Then we have $a_0 a_1 a_2 \dots a_n \leq (a_0 + a_1 + a_2 + \dots + a_n)^2 = m^2$.

This inequality is a special case of the Cauchy-Schwarz Inequality, which is a very fundamental and powerful inequality in mathematics. Here's a proof using Cauchy-Schwarz. Firstly, let's understand what the Cauchy-Schwarz Inequality states: For any sequences of real numbers x_1, x_2, \dots, x_n and y_1, y_2, \dots, y_n , the inequality is as follows:

$$(x_1 y_1 + x_2 y_2 + \dots + x_n y_n)^2 \leq (x_1^2 + x_2^2 + \dots + x_n^2) (y_1^2 + y_2^2 + \dots + y_n^2)$$

Now, let's apply this to prove the inequality you mentioned. Let $x_i = a_i$ and $y_i = 1$ for $i = 1, 2, \dots, n$. Then, the left-hand side of the Cauchy-Schwarz inequality becomes

$$(a_1 * 1 + a_2 * 1 + \dots + a_n * 1)^2 = (a_1 + a_2 + \dots + a_n)^2$$

and the right-hand side becomes

$$(a_1^2 + a_2^2 + \dots + a_n^2) (1^2 + 1^2 + \dots + 1^2) = (a_1^2 + a_2^2 + \dots + a_n^2) * n$$

Since the product $a_1 a_2 a_3 \dots a_n$ is one of the terms in the sum $a_1^2 + a_2^2 + \dots + a_n^2$ it follows that $(a_1^2 + a_2^2 + \dots + a_n^2) * n \geq a_1 a_2 a_3 \dots a_n$. Therefore, we have shown that

$$(a_1 + a_2 + \dots + a_n)^2 \geq a_1 a_2 a_3 \dots a_n$$

So, the inequality holds as per the Cauchy-Schwarz inequality. Please note, however, that this is a very specific case of the inequality where all a_i are non-negative. The Cauchy-Schwarz inequality holds more generally for any sequences of real numbers.

get-cheapest-option-in-time() takes $O(m^3)$ time. Because we have proven the number of all travel options is $O(m^2)$. For the worse case scenario where only the most expensive travel plan is within the given time constraints, we have to iterate through all $O(m^2)$ plan in which every plan can only involve $O(m)$ vertices since no loop is allowed. Then it will be $O(m^2(m)) = O(m^3)$ time.

In conclusion, the total time complexity is $O(m^3 + ml)$ time.

Solution 3.**(a)****High-level understanding**

We will first abstract a graph from the given n shapes. Every shape is a vertex in the graph. For any pair of shapes, we will examine if they (a) have an overlap and (b) have different number of sides. Only when these two conditions are met do we add this add an edge.

Then we adapt the Dijkstra's algorithm to find the shortest path from A to B . The difference is that we put weight on the vertices but not on the edges. But the spirit remains.

Detailed description

get-graph() To do this, first we abstract a shape into a vertex with the following quality: (1) number of sides, (2) area, (3) a list of vertices that it is connected to.

Then we iterate through all the shapes and find the edges that satisfy the condition that (a) they have an overlap and (b) they have different number of sides.

The upshot is that we have a graph that have all the possible vertices and possible edges.

Dijkstra-traverse() To traverse it, we slightly modify the Dijkstra's algorithm.

From vertex A in the remaining priority queue, we look at all the adjacent vertex K we first assign the value of $\text{Area}(A)$ to the edge (A, K) . After updating the edge, we continue with the Dijkstra's algorithm as usual.

After calling **get-graph()** first and **Dijkstra-traverse()** later, we have obtained the shortest path from A to all other vertices. And B is one of them, and thus the correctness is proven.

(b) The correctness is proven as the Dijkstra's algorithm is correct as in the lectures.

Claim 3.1. *The two filtering conditions in get-graph() is justified.*

Proof. It is as our algorithm uses it to ensure that all the edges in the graph are valid. \square

Claim 3.2. *The modified Dijkstra's algorithm is correct.*

Proof. Because we only have weights associated with vertices, the burden of proof lies in whether or not the the weight assigning process is justified. In this case, we need to show that the weight for each vertex after updating shows the minimum cost to get to that vertex.

We prove by induction.

Base case We are still at vertex A and the weight is the area of A itself.

Inductive hypothesis We have visited a set S of vertices and the weight attribute of the vertices in them have been updated correctly to reflect the costs to travel from vertex A to that visited vertex.

Inductive step To visit a new vertex K , for an edge between S_i and K where $S_i \in S$, we set the weight of edge (S_i, K) to $\text{Area}(K)$.

After that, when we are performing the *decrease-key()* operation, we notice that when we add the weight of edge (S_i, K) to the key of S_i , we have successfully obtain the value to get from A to K via S_i since the value of $\text{Area}(K)$ is the only difference between the current weight of paths and the previous one. \square

In conclusion, the two claims are proven and the rest is the same as the standard Dijkstra's algorithm as proven in the lectures.

(c) Building a graph from the shapes take $O(n^2)$ time and using the modified Dijkstra's algorithm takes $O(n^2)$ time. In total, $O(n^2)$ time.

For building the graph *get-graph()*, we have n shapes in total. For each shape K , we look at all other $n - 1$ shapes and conduct two $O(1)$ operations to examine if (a) they have an overlap and (b) they have different number of sides. That is $O(n)$ time for each shape, which yields $O(n^2)$ time for all n shapes.

For the modified Dijkstra's algorithm *Dijkstra-traverse()*, we have n vertices and at most $n(n - 1)/2$ edges. We do one more assigning value operation which takes $O(1)$ time. As proved in the lectures, using the Fibonacci heap, the algorithm takes $O(m + n \log n) < O(m + n^2)$ time. In our case, number of edges $m = O(n^2)$ because $m \leq n(n - 1)/2$, so the algorithm takes $O(n^2)$ time.

Combining the two parts, we have $O(n^2)$ time in total.